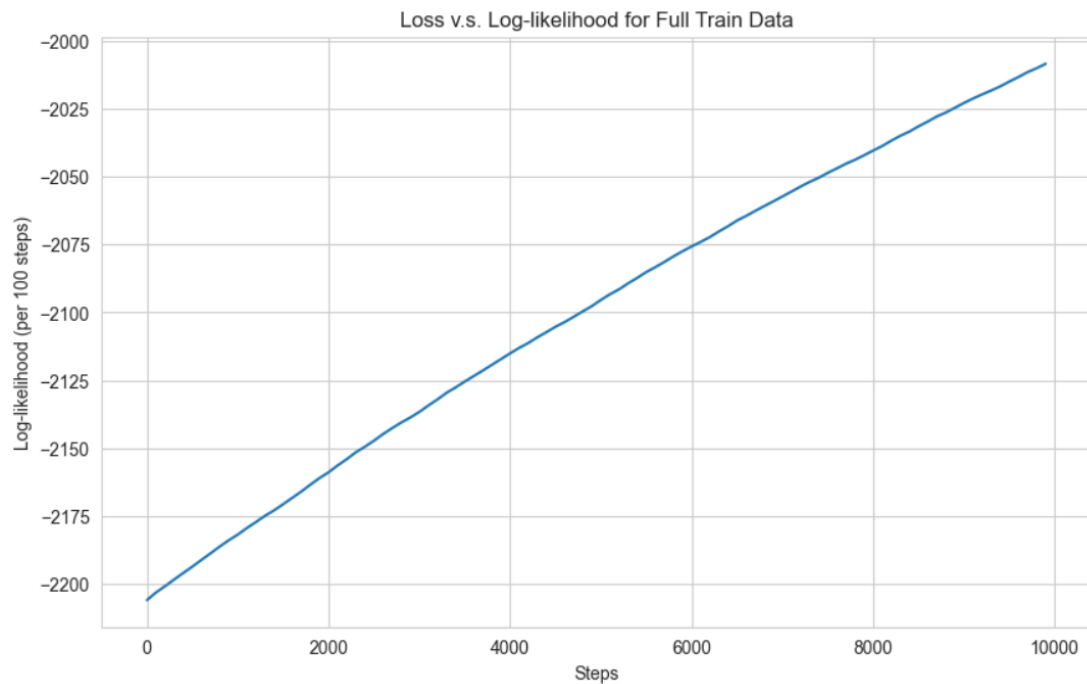
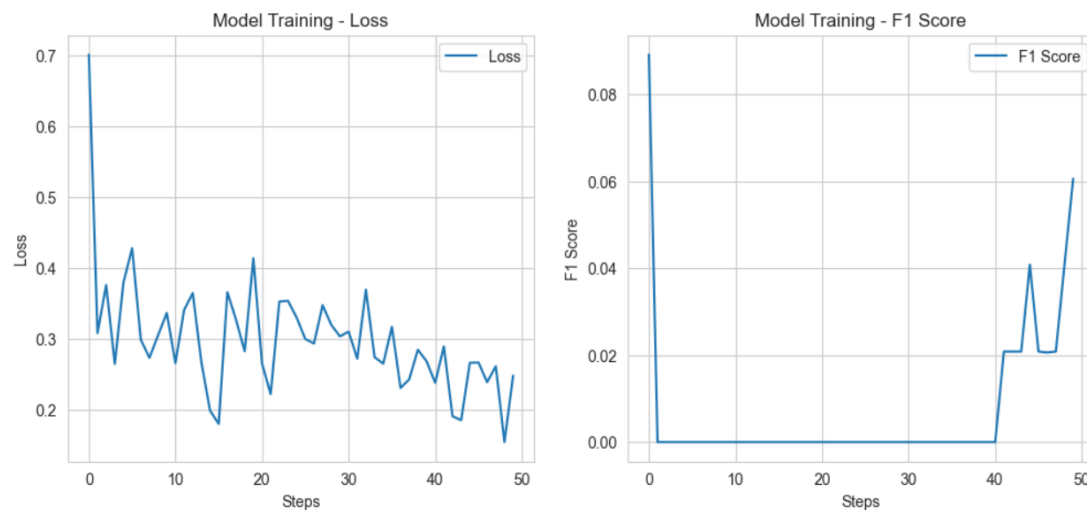


**Username on Kaggle: yanzhuo2001**


## **Part 2.1: Plot log-likelihood**



## **Part 3.1: 1000 steps training and reporting losses every 20 steps**

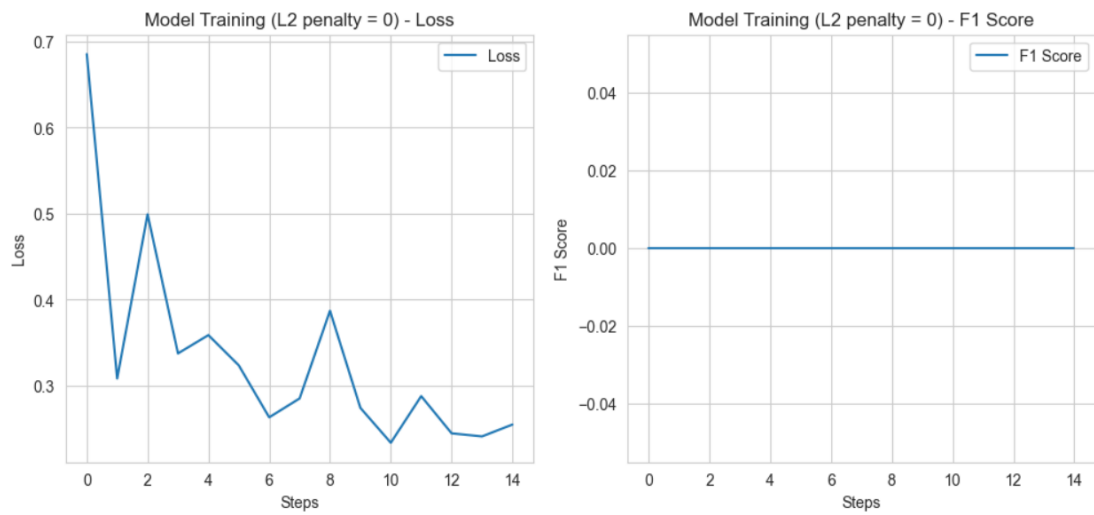


## **Part 3.2: Train at least 1 epoch and calculate the loss and F1 score**

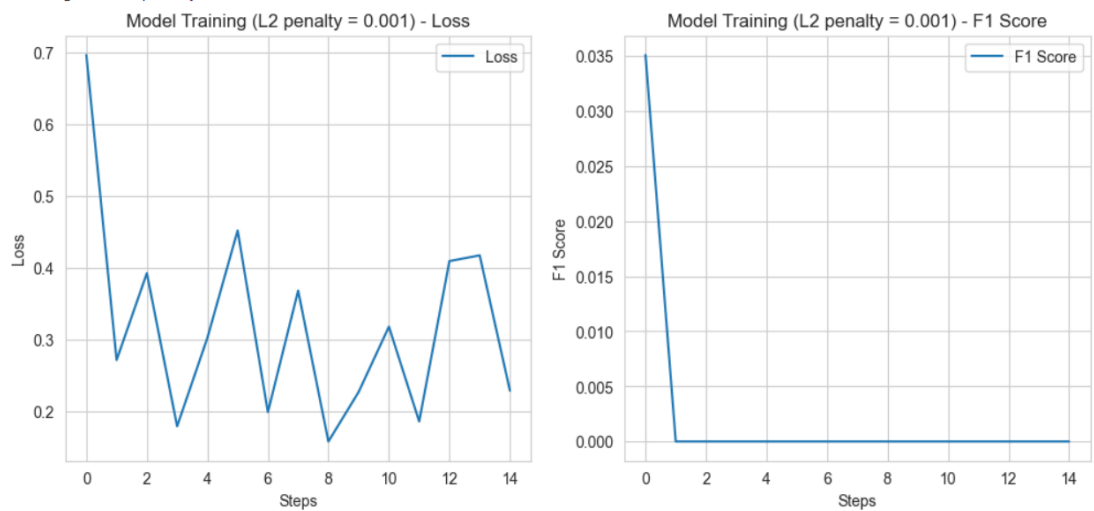


### Part 3.3: Add regularization (L2 penalty)



Training with L2 penalty = 0



Training with L2 penalty = 0.001

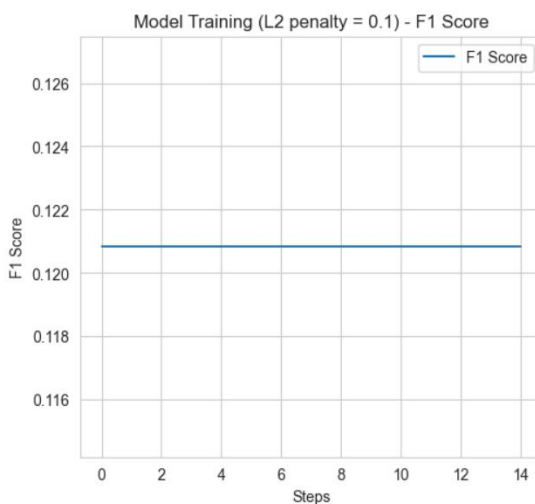



Training with L2 penalty = 0.1





## Part 3.4: Use different optimizers

Training with optimizer RMSprop

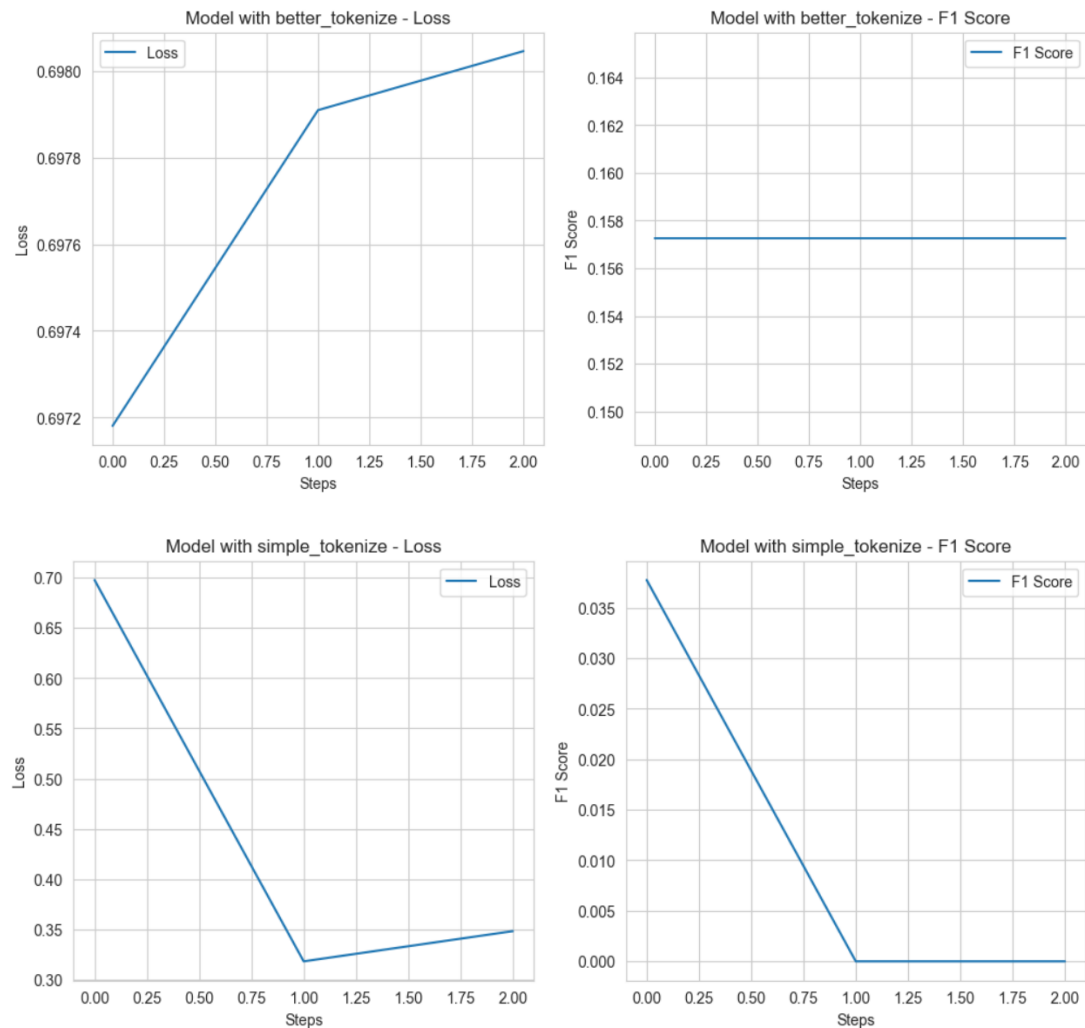


Training with optimizer AdamW



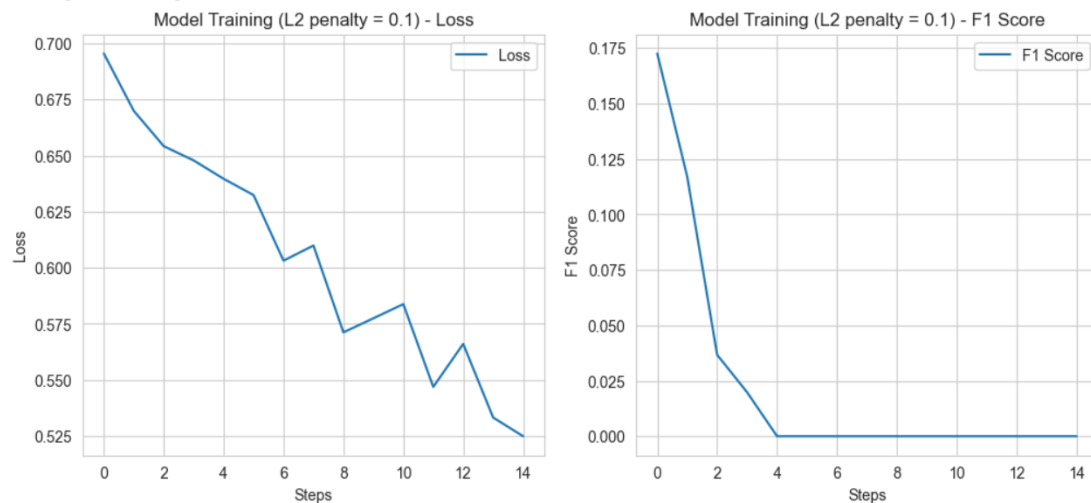
## Part 3.5: Comparison of different word segmentation methods

Epoch 0, Step 0, Loss: 0.6971805691719055, F1 Score: 0.1572  
Epoch 0, Step 50, Loss: 0.6979093551635742, F1 Score: 0.1572  
Epoch 0, Step 100, Loss: 0.6980456113815308, F1 Score: 0.1572


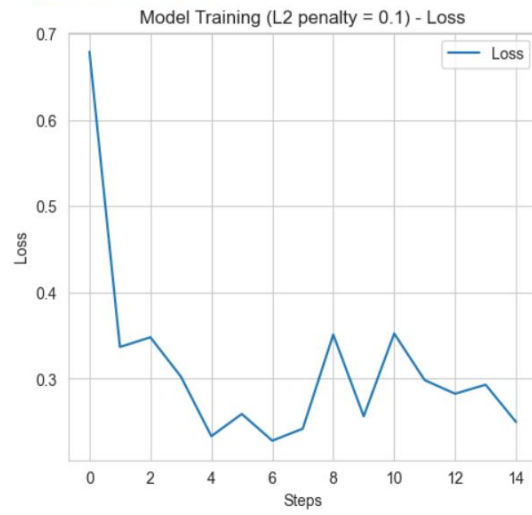


## Part 3.6: The effect of learning rate

Training with learning rate 0.001



Training with learning rate 0.1



Training with learning rate 1

